

Question 1: Indexing

1. List the names, ages, and offices of professors of a user-specified sex (male or female) who have a user-specified research specialty (e.g., artificial intelligence). Assume that the university has a diverse set of faculty members, making it very uncommon for more than a few professors to have the same research specialty.
 - Attributes: `sex`, `specialty`
 - Index Type: Composite index on (`sex`, `specialty`)
 - Unclustered, updates are not frequent
 - Index structure: B+ tree (range queries)
 - o **Query:** CREATE INDEX idx_prof_sex_specialty ON Prof(sex, specialty);
2. List all the department information for departments with professors in a user specified age range.
 - Attributes: `age`, `dept_did`
 - Index Type: Index on `age`
 - It is Unclustered
 - Index structure: B++ tree
 - o **Query:** CREATE INDEX idx_prof_age ON Prof(age);
3. List the department id, department name, and chairperson name for departments with a user specified age range.
 - Attributes: `num_majors`, `did`, `dname`, `chair_ssno`
 - Index Type: Composite index on `num_majors`
 - It is Unclustered
 - Index structure: B++ tree
 - o **Query:** CREATE INDEX idx_dept_num_majors ON Dept(num_majors);
4. List the lowest budget for a department in the university.
 - Attributes: `budget`
 - Index Type: Index on `budget`
 - It is Unclustered
 - Index structure: B++ tree
 - o **Query:** CREATE INDEX idx_dept_budget ON Dept(budget);
5. List all the information about professors who are department chairpersons.
 - Attributes: `chair_ssno`
 - Index Type: Index on `chair_ssno`
 - It is Unclustered
 - Index structure: B++ tree
 - o **Query:** CREATE INDEX idx_dept_chair_ssno ON Dept(chair_ssno);

Question 2: Storage and Indexing

Given

- Relation: `Student(sid, sname, major, email)`
- `sid` is the key
- `sid` values: uniformly distributed between `100` and `204,900`
- Attributes: `char(40)`
- 100,000 records
- Block size: 16KB + 8 bytes
- Record pointer size: 8 bytes

(a) Heap File Costs

a. File Scan:

- # of records per block = $\frac{16 \times 1024}{40 \times 4} = 1024$
- # of blocks = $\frac{100,000}{1024} = 98$
- Cost = 98D

b. Equality Search (`sid` = `25700`)

- Average case: Half the blocks need to be scanned
- Cost = $\frac{98}{2} = 49D$

c. Range Search (`sid` ≤ `25700`):

- Selectivity = $\frac{25700-100}{20480} = 0.125$
- Number of records to scan = $0.125 \times 100,000 = 12,500$
- # of blocks to scan = $\frac{12,500}{1024} \approx 13$
- Cost = 13D

(b) Clustered B++ Tree Index Costs

1. File Scan:

- # of blocks with 67% occupancy = $1.5 \times 98 \approx 147$
- Cost = 147D

2. Equality Search (`sid` = `25700`):

- Height of B+ tree = 3
- Cost = $Height + 1 \text{ for data} = 3 + 1 = 4D$

3. Range Search (`sid` ≤ `25700`):

- Selectivity = .125
- Cost for leaf pages = 13D
- Total cost = 4D (for tree) + 13D (for data) = 17D

(c) Unclustered B+ Tree Index Costs

1. File Scan:

- Same as heap file = 98D

2. Equality Search(`sid`= `25700`):

- Height of B+ tree = 3
- Cost = Height + 1 for data = 3 + 1 = 4D

3. Range Search (`sid` <= `25700`):

- Selectivity = 0.125
- # of records to scan = 12,500
- Each look up needs 1 I/O for index + 1 I/O for data
- Total cost = 12,500 I/Os (assuming random I/O)